



# Real models as points — descriptive placement (NOT a trilemma proof)

Panel 1 — per-model reliability (lead)



Panel 2 — (H, C) placement; A encoded as marker size



Descriptive placement of real, named open-weight models on the behavioral axes; this is NOT a trilemma proof, an impossibility result, or a traced trade-off region. The autonomy axis A is approximately 1.0 for every model because the competence probe ran ungated (no abstention incentive / no gating reward), so A is pinned by construction and the apparent absence of an autonomy trade-off is an artifact of the design, not evidence. A is itself only a proxy: it is the answer-commitment rate, where a task counts as 'abstained' whenever no parseable answer was produced -- which conflates deliberate deferral with mere unparseable-output (parse-fail) responses, so A measures behavioral answer-commitment, not autonomy-as-chosen-deferral. Consequently the (H, C, A) panel is competence-confounded and cannot be read as a causal trade-off: model placement there reflects task competence, not a mechanism. The causal trilemma claim rests on the gated mechanism experiment (hypotheses H1/H2/H4/H5) together with the theory, not on this scatter. What the calibration panel does show is each model's per-output reliability structure (reported confidence  $r$  versus realized correctness), i.e. the direction and magnitude of per-model overconfidence.